

NIFTY-50 Investment Intelligence Platform

Submitted By:

Roman Kumar

Enrollment No: 23117121

Branch: Mechanical Engineering

Year: 4th Year

1. Introduction & Project Objective

The financial markets generate large volumes of complex data daily. For an investor, extracting clear patterns from raw prices can be difficult due to market noise, structural trends, and changing volatility regimes. The primary purpose of this project is to build an end-to-end investment intelligence platform that processes historical market records and turns them into clean, structured, and actionable investment decisions.

By focusing on the major corporate symbols listed within the NIFTY-50 index, this platform addresses every critical stage of quantitative wealth management. This includes parsing and cleaning historical data records, generating high-utility technical alpha features, utilizing machine learning classification models to discover price direction clues, assessing individual stock risk cross-sectionally, and deploying multi-asset portfolio logic optimized for distinct investor profiles. The goal is to replace human bias with an automated, explainable framework built entirely on data evidence.

2. Exploratory Data Analysis (EDA)

The foundational analysis relies on a long-term database tracking daily closing figures, trading volumes, and value turnovers from January 2000 to April 2021. This specific timeline is highly valuable because it exposes our models to multiple major historical macro cycles, including the 2008 global credit crisis, the 2015 emerging market corrections, and the extreme liquidity fluctuations observed during the 2020 pandemic timeline.

Initial structural data checks revealed significant discrepancies in exchange reporting structures over time. For instance, detailed transaction column parameters like 'Trades' were missing close to 48.8% of their entire historical timeline. This happened because exchanges did not mandate trade count reporting until the latter half of the data collection cycle. To maintain statistical consistency across columns without introducing artificial imputation flaws, columns suffering from large systematic missing intervals were systematically filtered out. Chronological continuity was explicitly enforced across all stock sub-frames by parsing the date fields into uniform structures and sorting records by stock symbols and time coordinates.




Figure 1: Top 10 Companies by Number of Records. This graph details the active history available across leading blue-chip entities. Stocks like ASIANPAINT, SUNPHARMA, GRASIM, and CIPLA hold over 5,300 continuous data rows each, establishing a dense historical environment for validation testing.

A major risk when building automated quantitative systems across a large group of stocks is nominal scale variation. Stocks in the index trade at widely varying absolute values, often driven by corporate share split histories rather than intrinsic performance differences. To highlight this

structural property, the average nominal closing values of the leading stock classes were calculated and compared.




Figure 2: Top 10 Stocks by Average Closing Price. This visualization illustrates the severe scale differences between stock prices. High nominal value equities like SHREECEM trade above an average of 18,000 INR, while other high-performing stocks like INFY operate under 1,000 INR. This proves that raw price scales are non-stationary and cannot be fed directly into standard models.

To further understand long-term stock trajectories before testing predictive models, we calculated the total cumulative returns from the absolute origin point of the dataset to its termination in 2021. This step highlights the wide variation in performance across different asset classes over a multi-decade timeline.




Figure 3: Top 10 Stocks by Cumulative Historical Returns. This graph illustrates the immense variance in wealth compounding capacity over the 21-year timeline. Certain outliers like SHREECEM achieved substantial long-term growth, while other blue-chip companies maintained steady, lower-velocity performance.

3. Advanced Technical Feature Engineering

To extract meaningful signals from raw open, high, low, close, and volume sequences, multiple technical features were synthesized. These calculations were applied group-wise across individual stock partitions to prevent cross-sectional data leakage. This process transforms absolute price levels into scale-free indicators that track trends, momentum, and volatility boundaries:

- **Price Range:** This feature measures the absolute intraday spread between the daily high and daily low prices. It captures within-session price dispersion, showing the daily trading variance independent of the stock's absolute price level.
- **Intraday Price Gap:** This attribute checks the difference between the current session's open price and the previous session's close price. It captures overnight sentiment adjustments and morning auction order imbalances.
- **Daily Return:** Computes the percentage change in closing prices between consecutive trading sessions. This transforms non-stationary absolute prices into stationary, scale-free return distributions.
- **Simple Moving Average (SMA20):** Computes the clean arithmetic mean of a stock's closing price over a rolling 20-day window, smoothing out random daily noise to reveal the underlying intermediate trend.
- **Exponential Moving Average (EMA20):** Similar to SMA20, but uses an exponential decay weighting scheme to place more emphasis on immediate price changes. This helps the indicator react more quickly to sudden trend changes.
- **Relative Strength Index (RSI):** A classic momentum oscillator bounded on a strict scale from 0 to 100. By evaluating the ratio of exponential average gains to average losses over a 14-day trailing window, it flags whether an asset is entering overbought or oversold conditions.
- **Moving Average Convergence Divergence (MACD):** Computes the difference between a fast 12-day EMA and a slow 26-day EMA. This calculation helps identify shifts in momentum direction and trend strength.
- **Bollinger Bands (High / Low Envelopes):** Plots dynamic volatility boundaries located exactly two rolling standard deviations above and below the 20-day SMA line. These bands expand during periods of high market stress and contract during consolidated ranges, providing a clear map of local price extremes.

Because rolling calculations require a minimum number of lookback days, they naturally generate blank or null rows at the beginning of each stock's historical sequence. To fix this and prevent

model failure, a strict cleansing operation was applied to remove these truncation artifacts, ensuring a continuous data matrix before starting the machine learning phase.

4. Predictive Modeling Methodology

Predicting financial markets requires careful handling of data to avoid backtesting illusions. A common error in standard machine learning is using random cross-validation splits, which can leak future information into past predictions. To prevent this forward-looking bias, this project uses a strict chronological timeline split: for each stock, the first 80% of historical transactions are allocated exclusively for training, while the remaining 20% are held back as an out-of-sample testing set to simulate live trading performance.

The platform configures two distinct analytical tasks to evaluate price movements:

1. **Classification Task:** The model is trained to predict a binary direction target. If the next day's return is positive, it is flagged as 1; if negative or flat, it is flagged as 0. This gives a directional baseline for short-term trend discovery.
2. **Regression Task:** The model maps features directly to continuous future return percentage values. This continuous signal provides the necessary inputs for running portfolio allocation models.

5. Model Architecture & Performance Diagnostics

The core modeling engine uses a parallelized **Random Forest Ensemble** architecture containing 100 independent decision trees. This ensemble methodology is highly effective at identifying non-linear relationships across technical features without requiring complex data scaling or variable transformations.

When evaluated over the extensive out-of-sample testing partition (comprising 46,714 individual session records), the directional classification model achieved a baseline accuracy score of 50.61%. While this score appears close to random probability, it accurately reflects the low signal-to-noise ratio typical of short-term financial data. Even a small, consistent statistical accuracy margin above 50% can be highly valuable when paired with rigorous risk management and broad asset diversification.

6. Portfolio Construction Logic

Once raw directional signals are generated, the platform passes them to a portfolio optimization module. Capital is distributed across the top 10 ranked assets to ensure adequate diversification and prevent concentration in any single stock:

- **Conservative Portfolio Profile:** Focuses on capital preservation and maximum insulation against market downturns. It screens for assets with below-average historical volatility and distributes capital using an inverse-volatility rule, giving the largest weight to the most stable stocks.
- **Balanced Portfolio Profile:** Targets a steady mix of growth and stability. It limits maximum asset volatility while weighting assets based on risk-adjusted return ratios, striking an institutional balance between risk and reward.
- **Aggressive Portfolio Profile:** Aimed at maximizing capital growth over the long term. It filters out any assets with negative risk-adjusted returns and ranks the remaining stocks by historical compounding performance, distributing capital using a growth-friendly structure.

7. Risk Assessment Methodology

To monitor portfolio safety and capture downside exposure, the platform tracks four core risk metrics across all assets:

- 1. **Asset Volatility:** Measures the standard deviation of daily returns, annualized by multiplying by the square root of 252, to track typical stock price fluctuations.
- 2. **Sharpe Ratio:** Evaluates the amount of excess return an investor receives per unit of total risk, assuming a normalized risk-free return rate of zero.
- 3. **Sortino Ratio:** Adjusts the Sharpe framework by focusing exclusively on harmful downside volatility, ignoring beneficial upside price spikes. This protects assets with strong upward momentum from being penalized.
- 4. **Maximum Drawdown:** Tracks the largest peak-to-trough decline in a stock's value across the historical timeline, highlighting the worst-case potential loss during major market corrections.

To help investors interpret these distinct metrics, the system scales and combines them into a single, cohesive **Composite Risk Score (scaled from 0 to 100)**. This provides an easy-to-understand baseline for comparing risk across different stock classes.




Figure 4: Top 10 Stocks by Annualized Volatility. This analysis reveals the wide range of risk behaviors across the index. High-beta financial and industrial stocks like UTIBANK and BHARTI exhibit annualized volatility levels above 45%, requiring careful risk-sizing to protect portfolios from large drawdowns.

8. Model Explainability

To prevent the machine learning models from operating as a "black box," the platform calculates feature importance scores by tracking the Mean Decrease in Impurity (MDI) across all decision nodes. This analysis shows exactly which technical indicators carry the most weight when making predictions.




Figure 5: Feature Importance Analysis. The model diagnostics show that transaction Volume and short-term Daily Returns have the highest predictive impact. In contrast, slower trend indicators like SMA20 or raw price layers provide lower immediate value.

9. Key Insights & Empirical Findings

Analyzing the long-term results from the notebook yields three major insights:

- **Risk-Adjusted Outperformance:** Tickers such as BHARTI, UTIBANK, and SHREECEM achieved the highest risk-adjusted efficiency scores over the out-of-sample data timeline.
- **Downturn Resilience:** Consumer staples like HINDUNILVR and NESTLEIND proved to be highly effective defensive assets, maintaining the lowest maximum drawdown metrics during severe market corrections.
- **Long-Term Compounding Power:** Capital growth was heavily dominated by key growth stocks, led by companies like SHREECEM, BAJAJFINSV, and ULTRACEMCO.

10. Conclusion

The NIFTY-50 Investment Intelligence Platform successfully demonstrates how unstructured market data can be transformed into clear, risk-managed portfolios. By replacing emotional bias with automated technical features and ensemble modeling, the platform offers a transparent and auditable approach to long-term wealth management. The use of multi-metric risk scoring and diversified capital allocation provides a robust framework that aligns investment decisions directly with data evidence.